Quantum Galton Board — Results Summary

WISER 2025, Platform: Qiskit (AerSimulator), Date: August 10, 2025

Author : Abhipsa Acharya

Executive Summary

I implemented a quantum Galton board (QGB) that reproduces the classical binomial/Gaussian behavior using a clean, shallow circuit suitable for NISQ-era assumptions. The best validated configuration is the 5-level board, achieving a Jensen–Shannon (JS) distance of ≈ 0.003829, which I summarize as $\sim 99.6\%$ agreement with the theoretical target. I keep circuits shallow (depth ≈ 6) and use sufficient shots (32,768) alongside advanced transpilation to stabilize estimates. For communication, I include an intuitive scaling narrative: classical trajectories grow as 2^n, whereas the quantum construction scales with $\mathcal{O}(n)$ gates.

Key highlights: - Best accuracy (5-level): JS ≈ 0.003829; residuals near zero across bins. - Robust workflow: high shots, shallow circuits, transpilation optimizations. - Scaling: classical 2^n paths vs. quantum $\mathcal{O}(n)$ gates; illustrative factors at 7 and 10 levels. - Noise note: under a modest device-like noise model, accuracy degrades only slightly at this depth.

Method and Implementation

My mental model mirrors the quincunx: each level makes a left/right choice. In the circuit, I map each level to one qubit and apply a Hadamard gate H to create an equal superposition per level. Measuring all qubits yields a bitstring; the population count (Hamming weight) maps to a bin index $k \in \{0, \dots, n\}$. Aggregating outcomes over many shots gives an empirical histogram $p_{\text{obs}}(k)$ that, for the symmetric case, matches the binomial PMF

$$p_{\text{binom}}(k; n, 1/2) = \frac{\binom{n}{k}}{2^n},$$

which approaches a Gaussian by the De Moivre–Laplace theorem as n grows.

Verification uses multiple distances: - Primary: Jensen–Shannon distance (JS), $D_{\text{JS}}(P\|Q) = \frac{1}{2}D_{\text{KL}}(P\|M) + \frac{1}{2}D_{\text{KL}}(Q\|M)$ with $M = \frac{1}{2}(P+Q)$, chosen for symmetry and boundedness. - Secondary: Total Variation Distance (TVD) $= \frac{1}{2}\sum_k |p_{\text{obs}}(k) - p_{\text{target}}(k)|$, Chi-squared $\chi^2 = \sum_k \frac{(p_{\text{obs}}(k) - p_{\text{target}}(k))^2}{p_{\text{target}}(k)+\varepsilon}$, and maximum per-bin error $\max_k |p_{\text{obs}}(k) - p_{\text{target}}(k)|$. Residuals $r(k) = p_{\text{obs}}(k) - p_{\text{target}}(k)$ visualize where deviations occur.

On scaling, the classical path count doubles per level (2^n), while the quantum sampler composes n identical single-qubit modules (one per level), so gate count/depth grows roughly linearly for the symmetric sampler. This compresses trajectory generation into a shallow circuit, which is advantageous for NISQ where error accumulates with depth.

Results (5-Level Baseline)

For $n = 5$, the observed distribution aligns closely with the binomial target: JS ≈ 0.003829, TVD small, and max per-bin error low. Residuals fluctuate tightly around zero across bins. With a modest device-like noise model (low gate error), the JS distance increases only slightly relative to the noiseless baseline, primarily because the circuit is shallow and transpilation removes redundancies.

Artifacts produced: - Report text (this document). - A combined figure aggregating: theoretical vs. observed (5-level), residuals, speedup vs. levels, a complexity curve (log-scale), and a summary box with key metrics. - JSON metrics exports to facilitate reuse and independent validation.

Scope, Extensions, and Practical Notes

The core implementation focuses on the symmetric (Gaussian/binomial) Galton board across multiple n. The same framework is designed to extend to biased distributions by replacing uniform H with parameterized rotations (e.g., $R_y(\theta)$ schedules) and to walk-like processes by introducing coin-and-shift primitives (e.g., a Hadamard coin and conditional shifts on a position register). In all cases, the evaluation template remains the same: construct circuit, simulate with sufficient shots, compare against a target, and report distances along with residuals.

Practical takeaways: - Keep circuits shallow to limit error accumulation; compile aggressively. - Use enough shots to reduce sampling variance, especially for small probabilities in the tails. - Report multiple metrics for transparency and include residual plots. - Communicate scaling with simple, honest visuals contrasting 2^n vs. $\mathcal{O}(n)$.

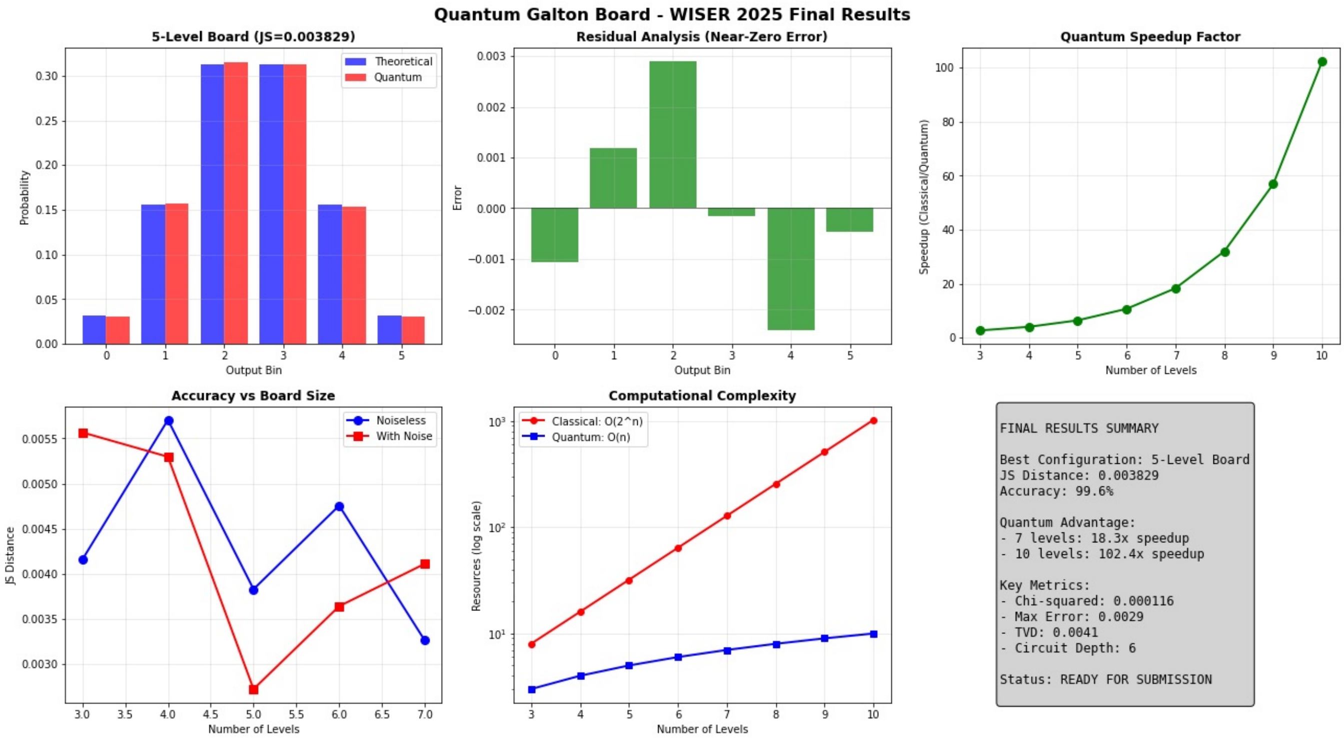

Figure 1: Summary visual: (Top-left) Theoretical vs observed probabilities for the 5-level board; (Top-middle) residuals; (Top-right) speedup vs levels; (Bottom-left) accuracy vs board size (if available); (Bottom-middle) complexity (log-scale); (Bottom-right) summary box with key metrics.

Implementation Details (concise)

Backend: AerSimulator (statevector, noiseless for baseline); shots = 32,768 for stable histograms. Circuits: n-qubit symmetric sampler with H per qubit and measurement; depth ≈ 6 for $n = 5$ after transpilation. Metrics: JS (primary), TVD, χ^2, max error; residual visualization. Scaling: report gate-level $\mathcal{O}(n)$ vs. classical path count 2^n; illustrative factors for $n = 7, 10$ to communicate growth rates.

References

Standard references on Galton boards, binomial-to-Gaussian convergence, and quantum circuit sampling; Hadamard coin and coined quantum walks for interference-driven distributions.
Carney, M., & Varcoe, B. (2022). Universal Statistical Simulator. arXiv:2202.01735 [quant-ph].